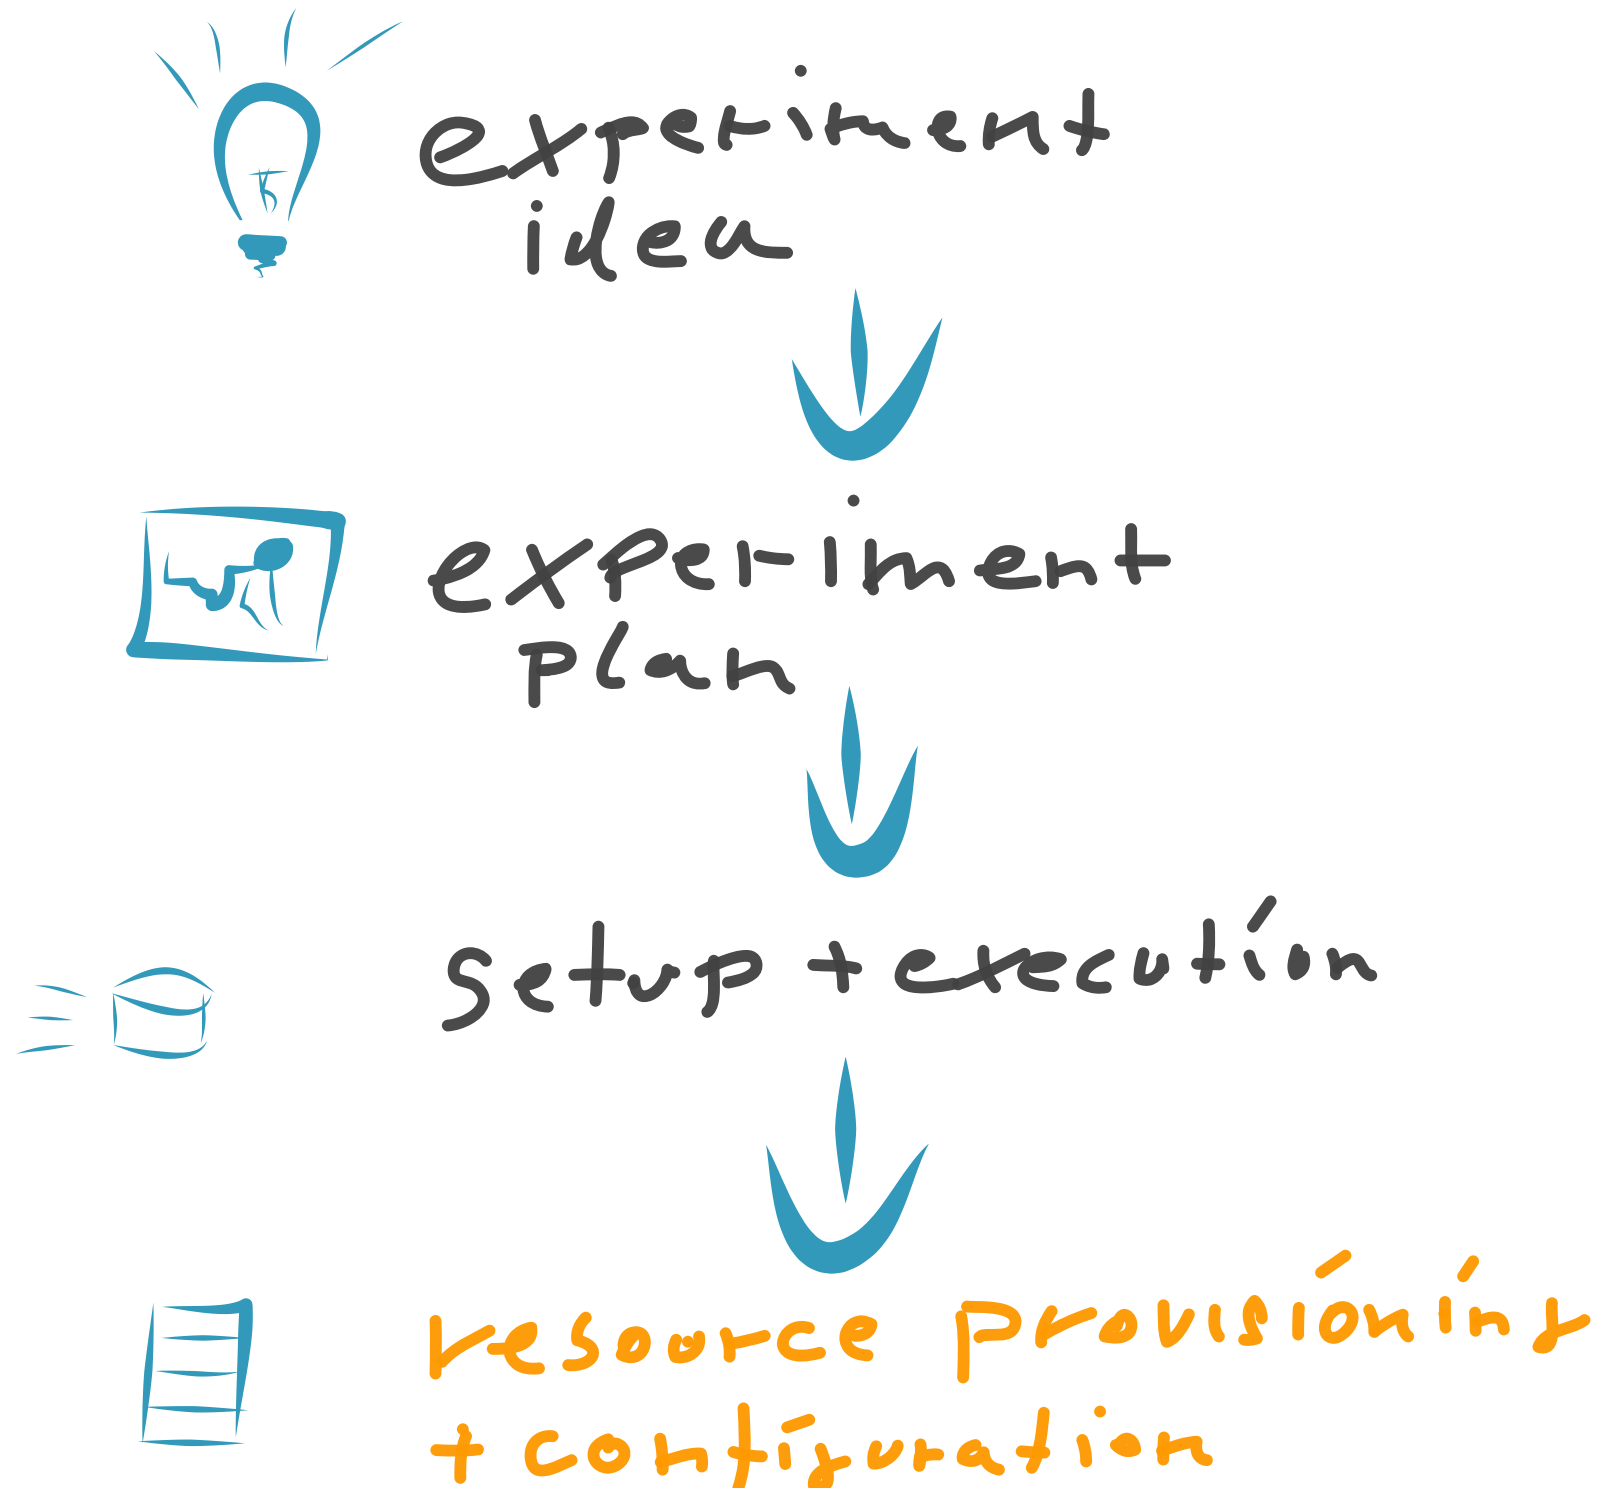


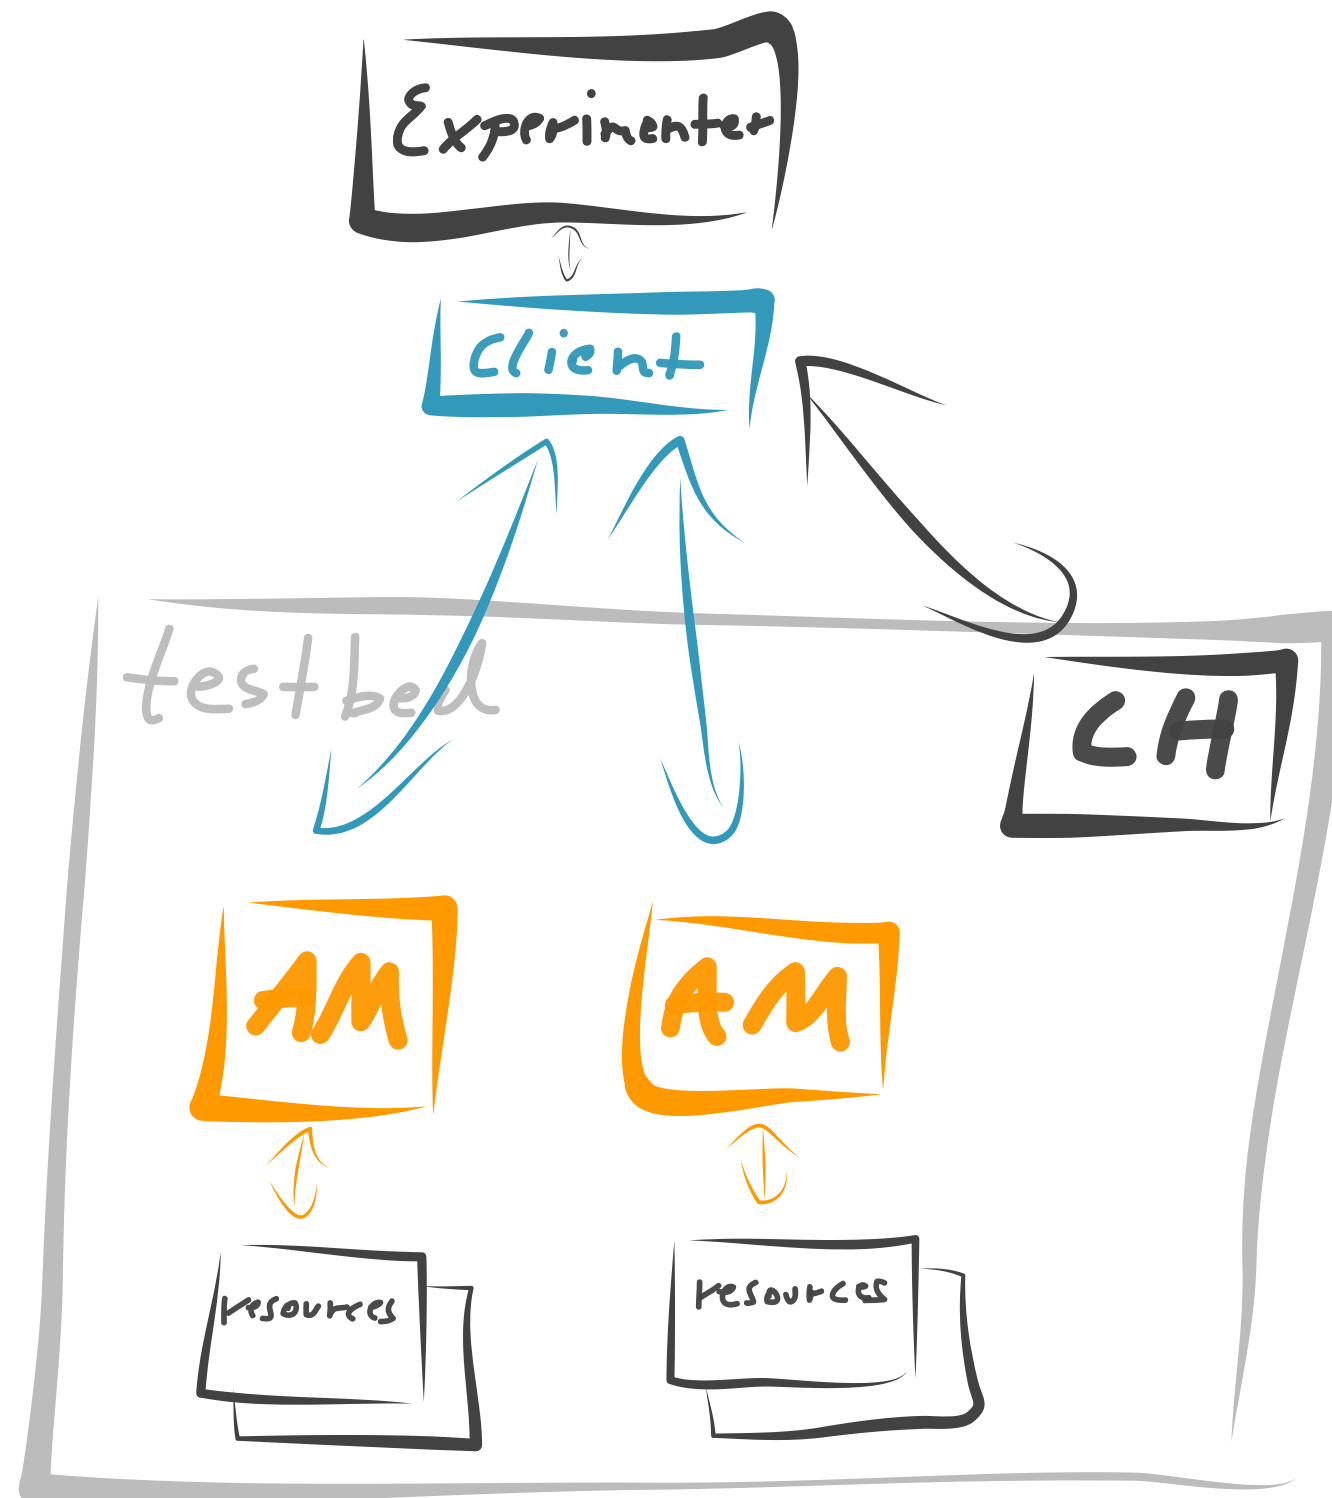
AMsoil

The glue for Aggregate Manager developers

researcher's goal




experiment execution



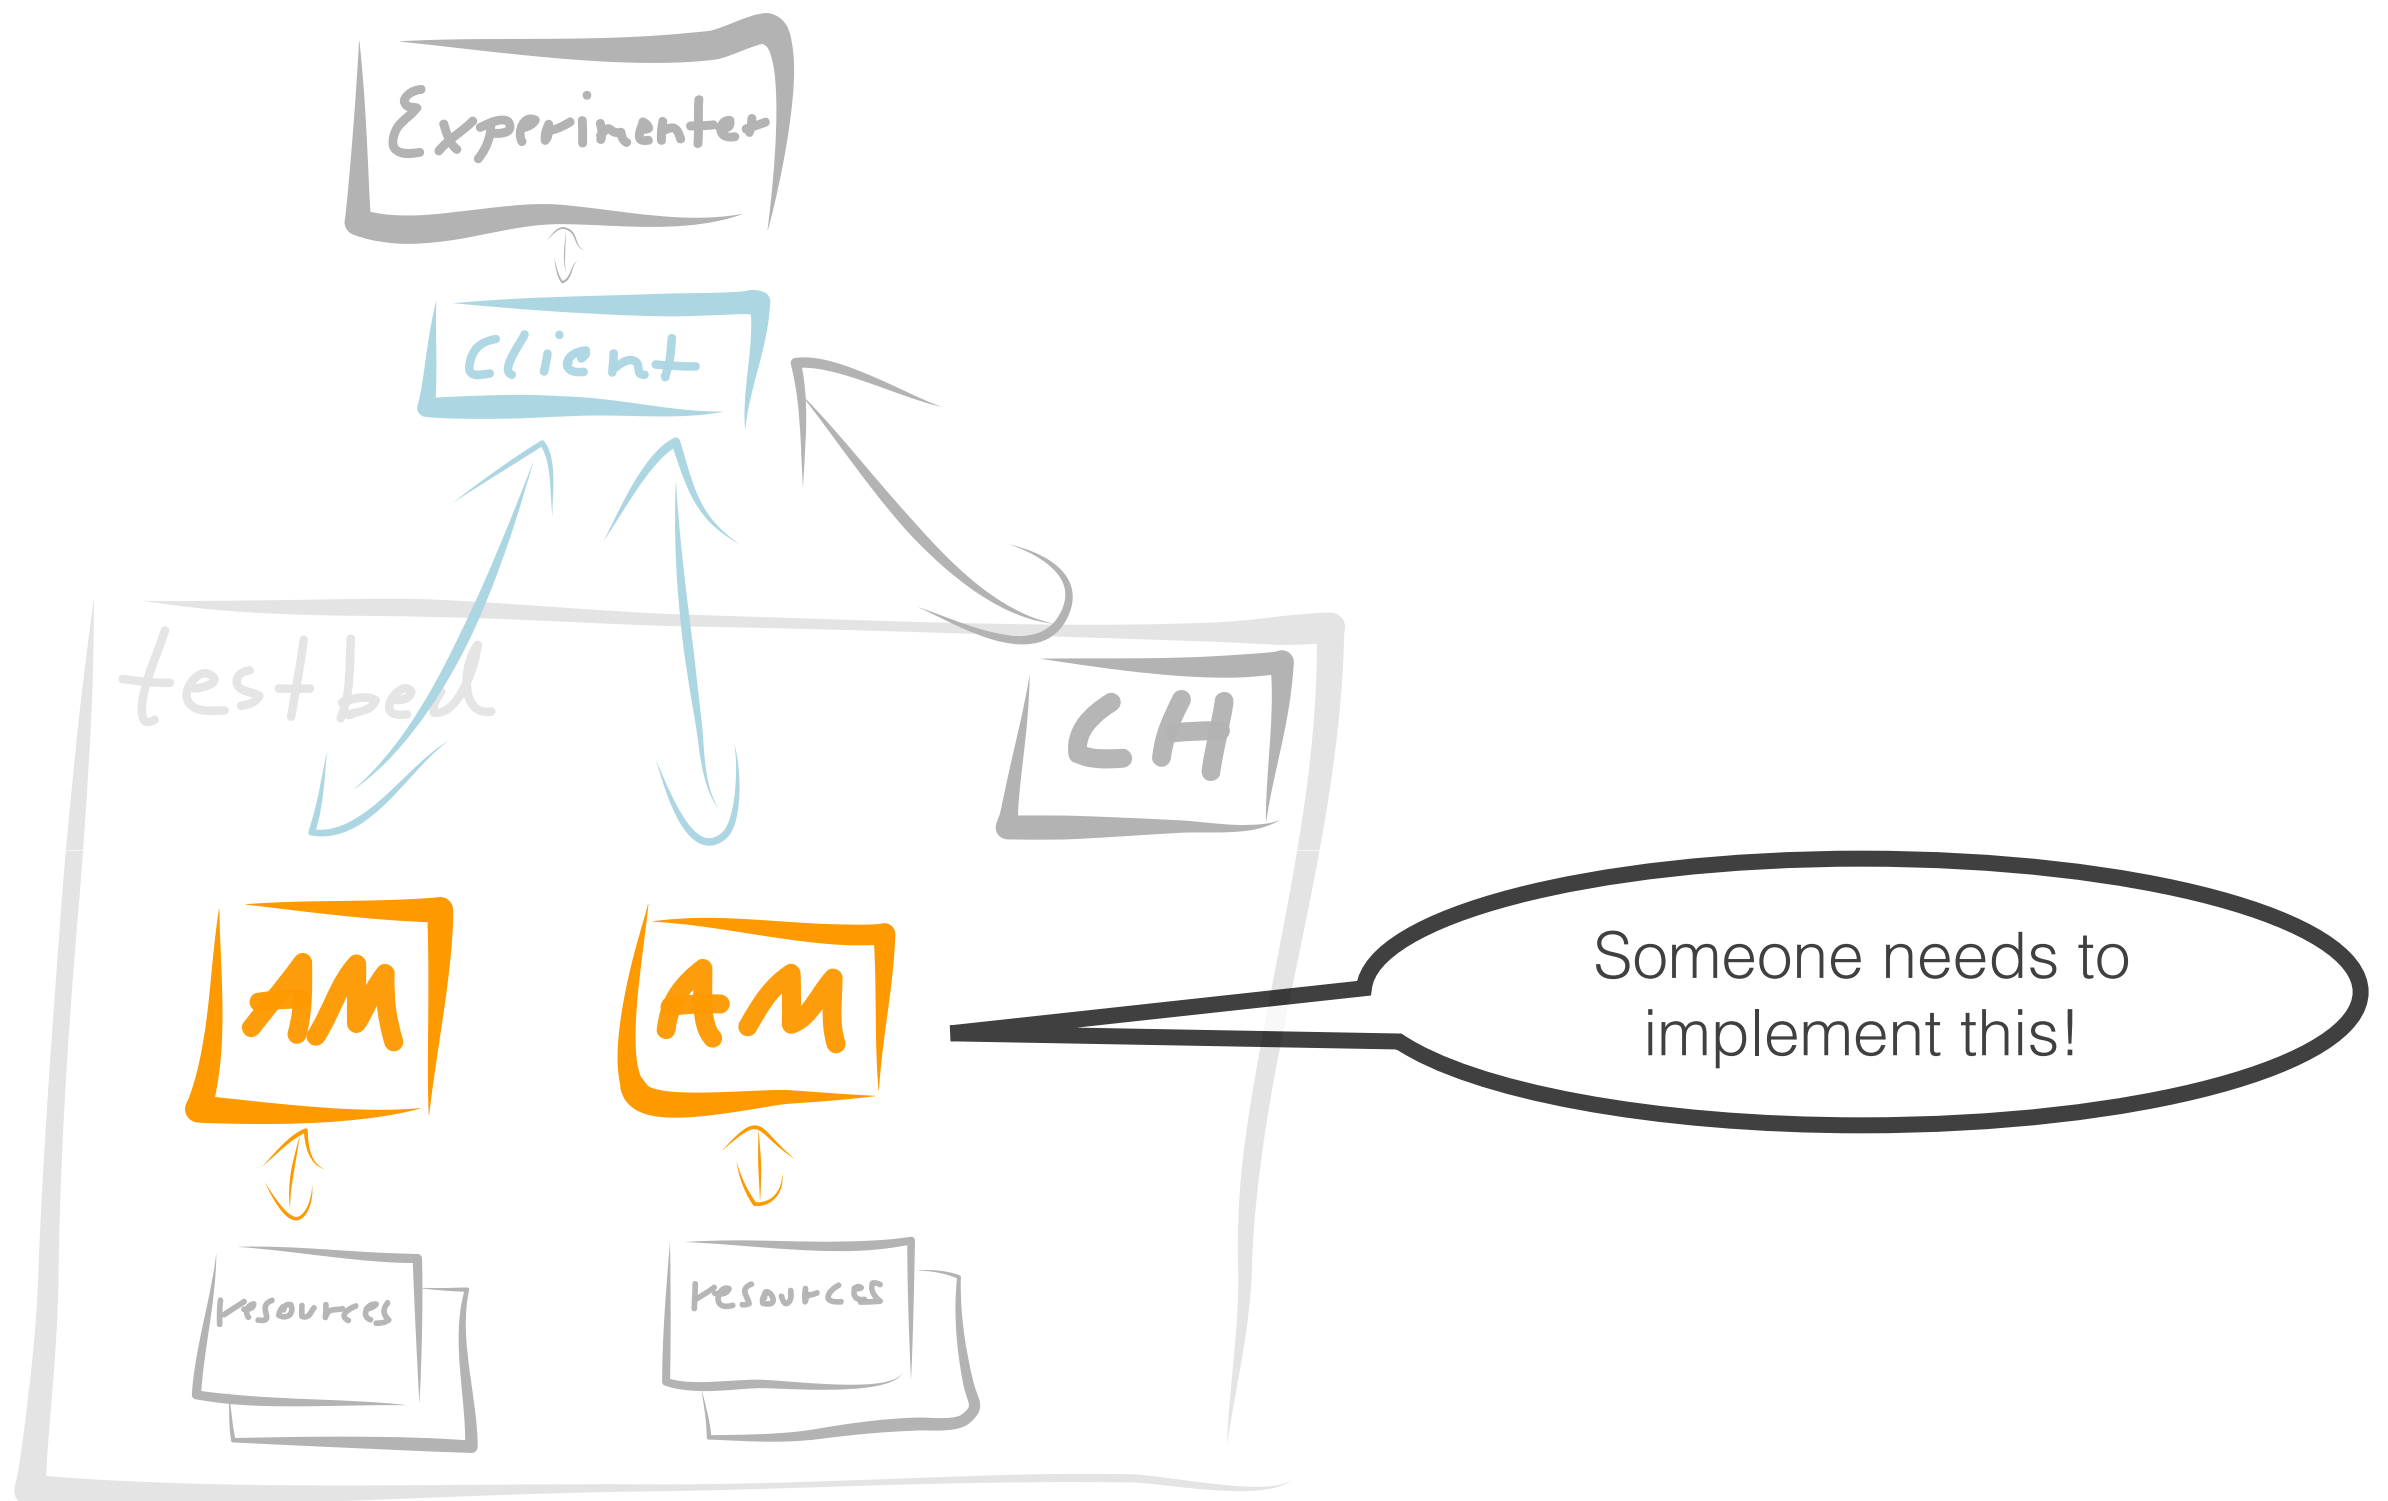
CH Clearinghouse
AM Aggregate Manager

test bed

- **Clearinghouse** manages certificates and credentials
- The **client** (*here: omni*) assembles the request and sends it to the Aggregate Manager
- **Aggregate Manager** manages, allocates and provisions resources



AMsoil?



AMsoil?




AMsoil is a light-weight [framework for creating Aggregate Managers](#) in test beds.

AMsoil is a pluggable system and provides the necessary glue between RPC-Handlers and Resource Managers . Also it provides helpers for common tasks in [AM development](#).

motivation


AM development
without AMsoil



This is why
you write an AM.
The rest is just annoying.


motivation

AM development
with AMsoil



how to write an AM

- Setup a little test bed
 - Install a Clearinghouse
 - Install a client
 - Install AMsoil
- Understand AMsoil
- Start hacking...



need to know

- how a **GENI** testbed works
- how **plugins** work
- what plugins you **need to develop**
- what else **AMsoil** supports

what now?

finish this presentation,

clone the repository [🔗 https://github.com/fp7-ofelia/AMsoil.git](https://github.com/fp7-ofelia/AMsoil.git)

then **read** [🔗 https://github.com/fp7-ofelia/AMsoil/wiki/Installation](https://github.com/fp7-ofelia/AMsoil/wiki/Installation)

GENI?

AMsoil managers are used in a GENI-like test bed.

Let's understand how GENI works.

names in GENI

- **Experimenter**

A human user who uses a client to manage resources via an AM.

- **Sliver**

A physical or virtual resource. It is the smallest entity which can be addressed by an AM

(e.g. an IP address, a virtual machine, a FlowSpace).

- **Slice**

A collection of slivers.




communication

- The Clearinghouse provides services to know who you are and what you may do.

(we don't care, just use it)

- The client speaks the GENI AM API to the AM.


(we care, because we implement it)



what can the API do?

<code>GetVersion</code>	Get info about the AM's
<code>ListResources</code>	Info what the AM has to offer
<code>Describe</code>	Info for a sliver
<code>Allocate</code>	Reserve a slice/sliver for a short time
<code>Renew</code>	Extend the usage of a slice/sliver
<code>Provision</code>	Provision a reservation for a longer time
<code>Status</code>	Get the status of a sliver
<code>PerformOperationalAction</code>	Change the operational state of a sliver
<code>Delete</code>	Remove a slice/sliver
<code>Shutdown</code>	Emergency stop a slice

allocate *and* provision?



allocated only for a short time resources are only booked not provisioned
provisioned the slice/sliver actually takes up resources (is actually usable)

typical experiment

Imagine a restaurant reservation.

- **ListResources**

Call the restaurant to ask what tables are available.

- **Allocate**

Call to tell which table you want (they will only hold the table for 2 hours).

- **Provision**


Come and use at the table (this may take 5 hours).

how do say what I want?

The resources are described with an XML document called RSpec.

There are three RSpec types:

- **Advertisement** (*short: ads*)
Announces which resources/slivers are available.
- **Request**
Specifies the wishes of the experimenter
- **Manifest**
Shows the status of a sliver



AM... what now?

Let's look on AMsoil and see what it can do.

a broad look

AMsoil's directory structure

```
|-- admin
|-- deploy
|   |-- trusted
|-- doc                                Documentation
|   |-- img
|   |-- wiki
|-- log                                AMsoil's log
|-- src
|   |-- amsoil                         AMsoil's core implementation
|       |-- core
|   |-- disabled-plugins              Unused code/plugins
|       |-- ...
|   |-- plugins                       Plugins to be loaded when bootstrapping AMsoil
|       |-- ...
|-- test
```

why plugins?

- **Selection**

An administrator can add/remove plugins/functionality.

- **Exchangeability**

The interface remains, but the implementation be changed.

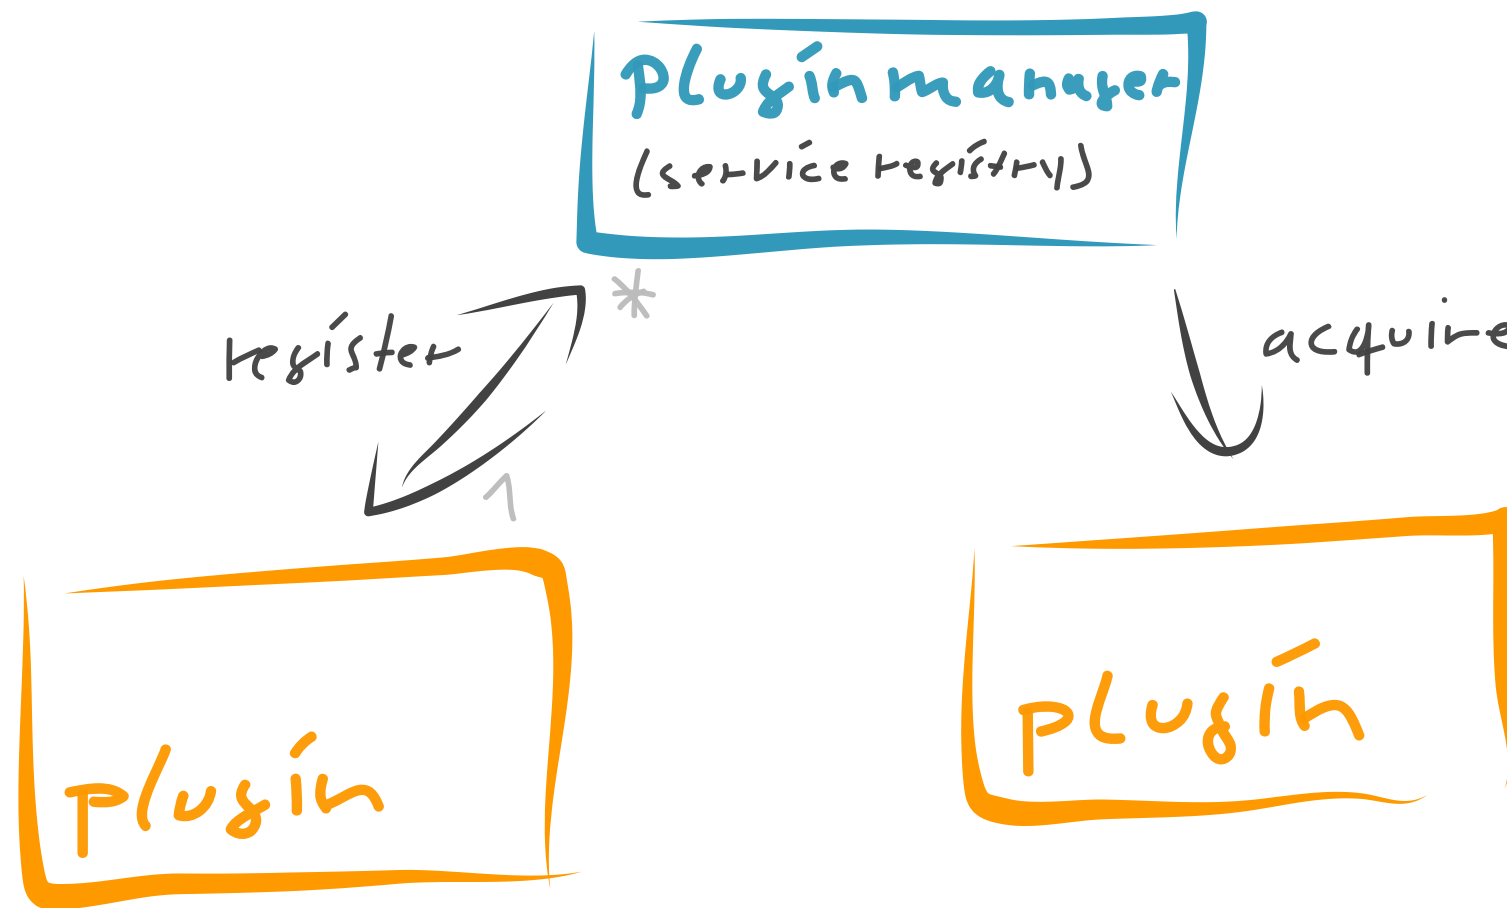
- **Clarity**

Provide a set of services and hide the details behind.

- **Encapsulation**

Protect implementations from other developers.

register and use plugins



```
[plugin A] import amsoil.core.pluginmanager as pm
[plugin A] pm.registerService('myservice', serviceObject)

[plugin B] service = pm.getService('worker')
[plugin B] service.do_something(123)
```

what can be a service?

short version

everything which can be referenced in Python

long version

ints, strings, lists, dicts, objects, classes, packages



yes even
packages!

under the hood

plugin

describes
services & dependencies

Manifest

performs
initialization & registration

plugin.py

implementation

implement a plugin

- create a [new folder](#) in plugins
- create the [manifest.json](#)
- create the [plugin.py](#)
 - write a [setup\(\)](#) method
- [register](#) your services

implement a plugin

manifest.json

```
{
  "name"          : "My Plugin Name",
  "author"        : "Tom Rothe",
  "author-email"  : "tom.rothe@eict.de",
  "version"       : 1,
  "implements"    : ["myservice", "myclass", "mypackage"], # you'll register these services
  "loads-after"   : ["somedependency"],                  # dependency needs to be loaded before the setup method
  "requires"      : []                                    # dependency can be loaded after the setup method
}
```

plugin.py

```
# ...imports...
def setup():
    # register a service
    pm.registerService('myclass', ServiceClass)
    pm.registerService('myinstance', SingleClass() )
    pm.registerService('mypackage', my.python.package)
```

@serviceinterface

The methods and attributes which can should be used are marked the annotation @serviceinterface.

implementation

```
from amsoil.core import serviceinterface

class MyService(object):
    @serviceinterface
    def do_something(self, param):          # can be used by the service user
        pass
    def do_more(self, param):              # not part of the service contract, NOT to be used
        pass
```

DOs and DON'Ts

- If you have plugin-specific exceptions, create a [package with all exceptions](#) and register the package as a service.
- Separate a plugin [into multiple plugins](#) if this improves re-usability.
- [Never import another plugin directly](#), always go via the pluginmanager via `pm.getService()`.

incoming missile

Let's find out how to react to RPC requests.

getting the requests

- **RPC Handler**

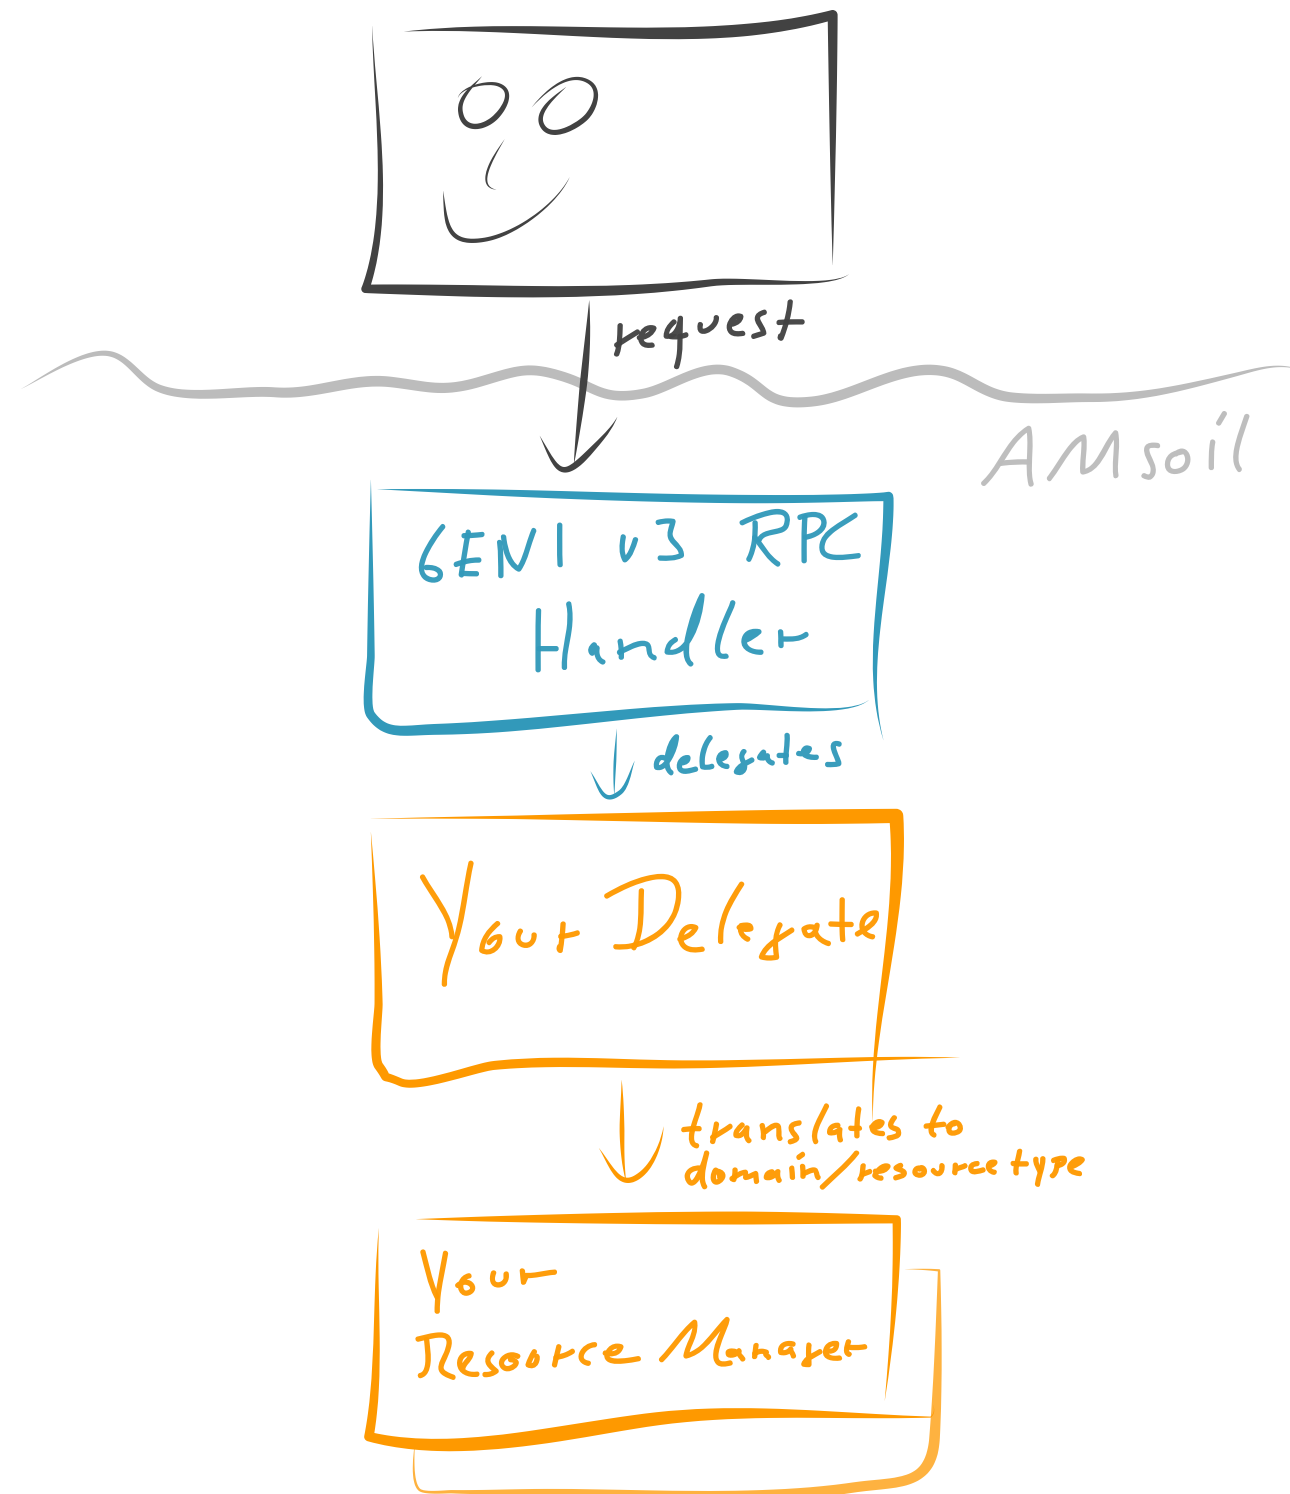
Retrieves the XML-RPC request, does some magic and passes the request on to the delegate.

- **Delegate**

Translates the GENI request into a language the Resource Manager can understand

- **Resource Manager** (short: RM)

Handles the actual allocation of the resources.



why RM *and* Delegate?



We need to **decouple** the RPC API from the resource management logic.

This enables AMsoil-based AMs to implement **multiple APIs** (e.g. GENI, SFA, OFELIA APIs) without having to re-write everything.

interfaces

- **Delegate**

Should derive from DelegateBase and overwrite the methods prescribed (e.g. list_resources, allocate, ...).

- **Resource Manager**

You make up the interface!

The methods, attributes, parameters are domain-specific and depend on the resource type being handled.

a new plugin is born

Create new plugins which handle the incoming requests from the client and do the actual resource management.

YourDelegate

- ✓ New folder for plugin
- ✓ manifest.json
- ✓ plugin.py
- ✓ a delegate object

YourResourceManager

- ✓ New folder for plugin
- ✓ manifest.json
- ✓ plugin.py
- ✓ a manager service

YourDelegate

yourdelegate/plugin.py

```
# ...imports...
GENIv3DelegateBase = pm.getService('geniv3delegatebase')
geni_ex = pm.getService('geniv3exceptions')

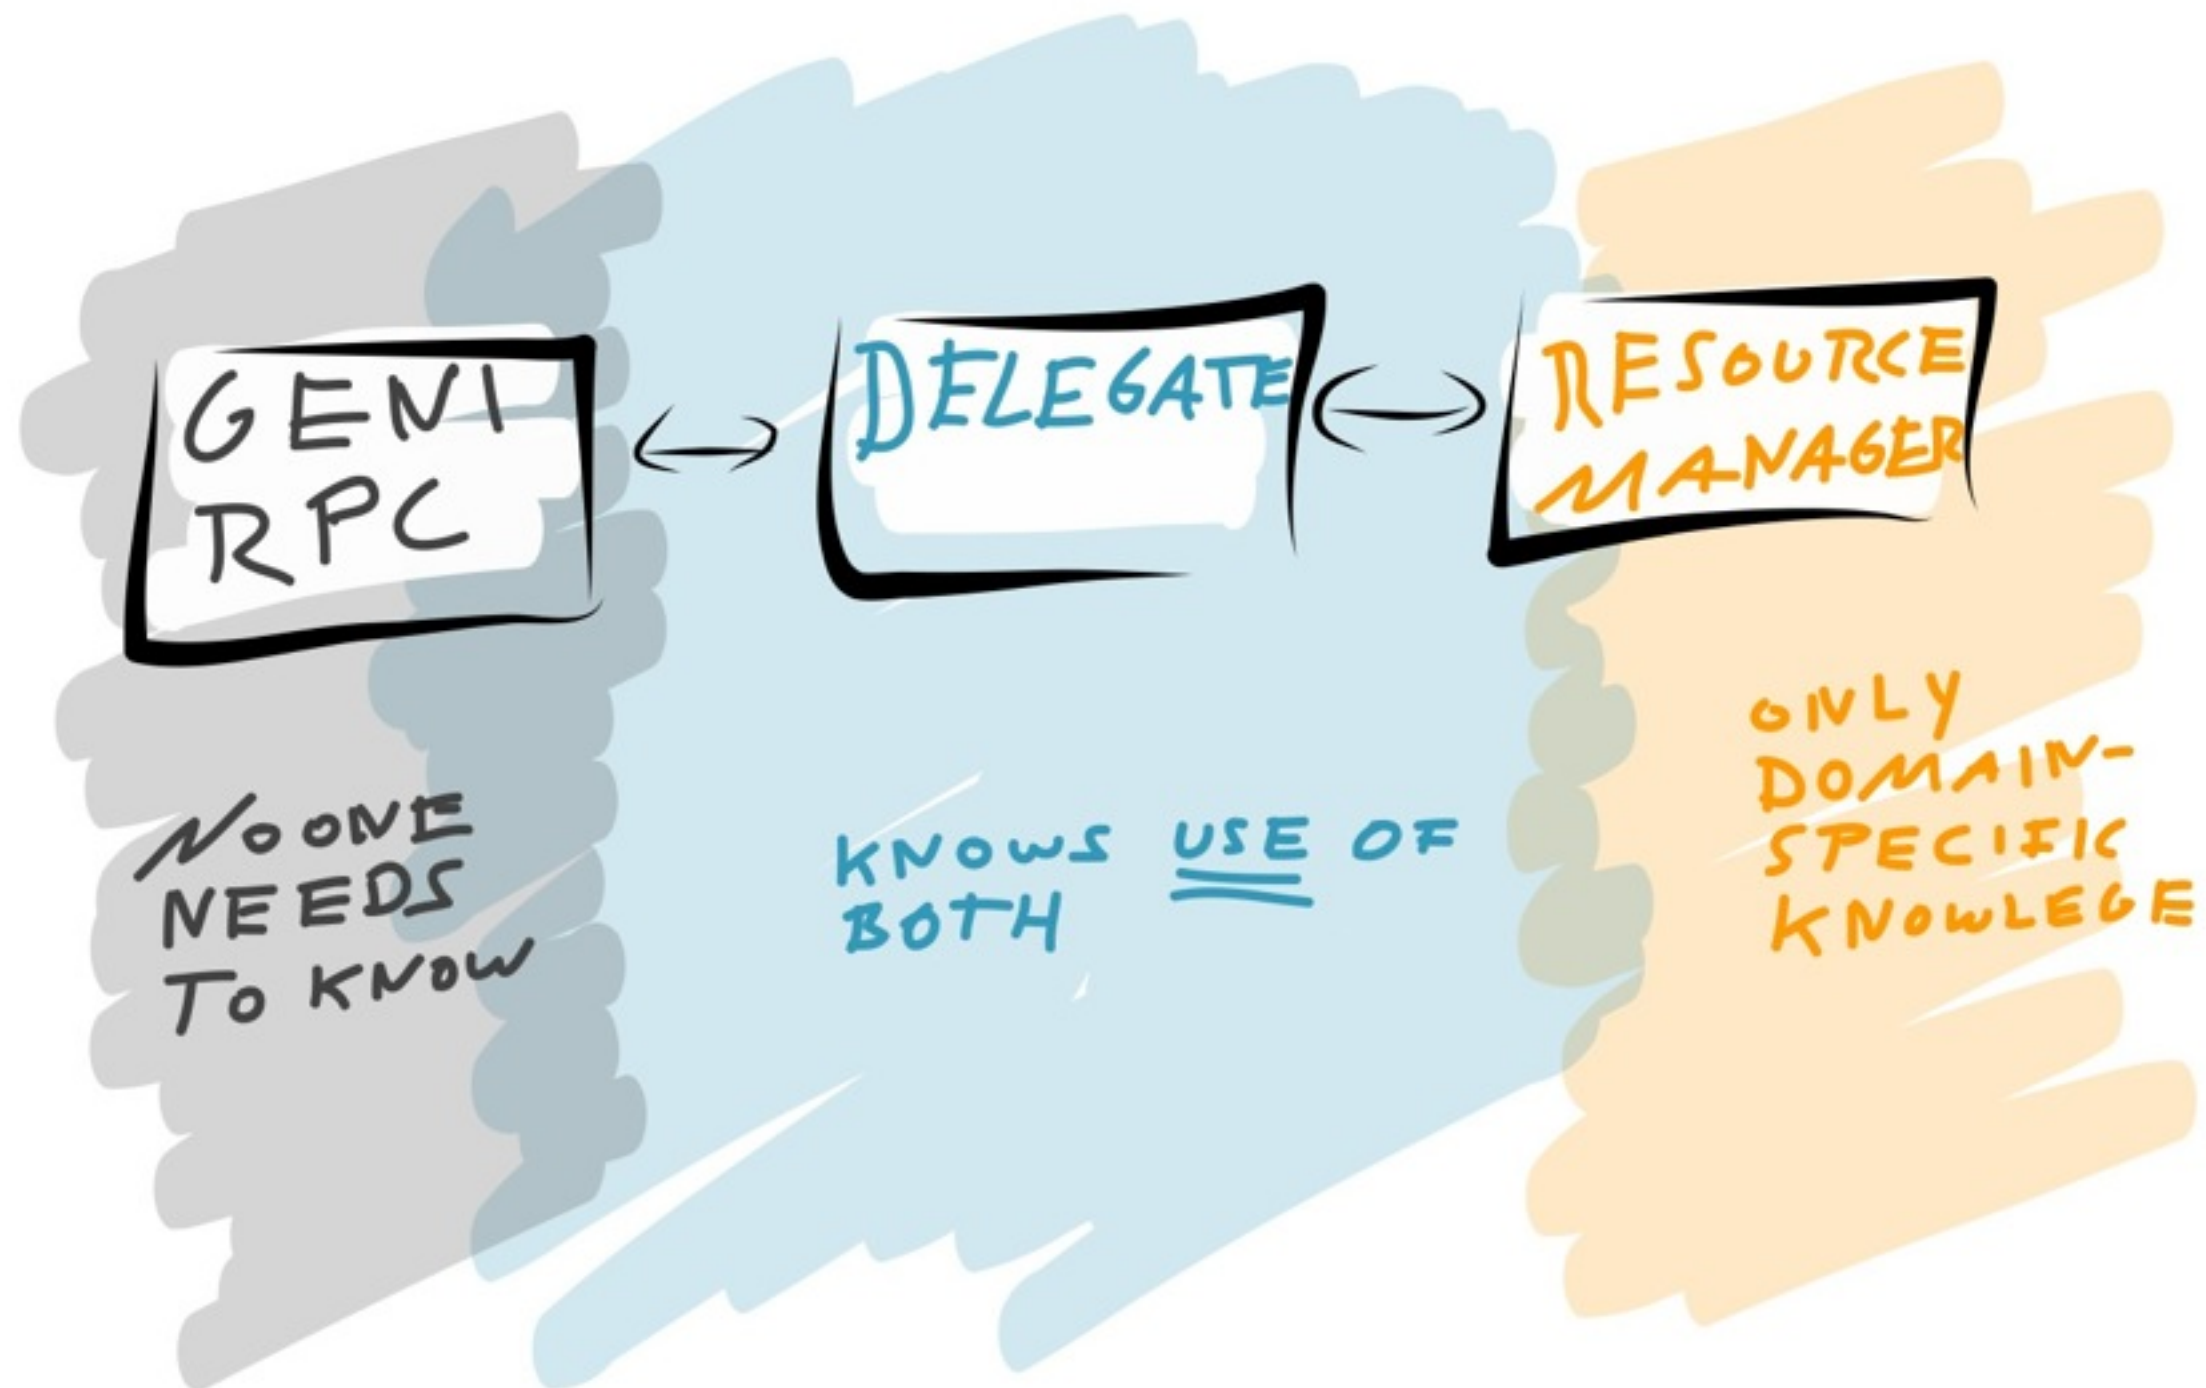
class MyDelegate(GENIv3DelegateBase): # derive from DelegateBase
    # ...
    def allocate(self, slice_urn, client_cert, credentials, rspec, end_time=None): # Overwrite DelegateBase method
        # perform authentication and check the privileges
        client_urn, client_uuid, client_email = self.auth(client_cert, credentials, slice_urn, ('createsliver',))

        rspec = self.lxml_parse_rspec(rspec) # call a helper method to parse the RSpec (incl. validation)
        # ...interpret the RSpec XML...
        try:
            # call a resource manager and make the allocation happen
            self._resource_manager.reserve_lease(id_from_rspec, slice_urn, client_uuid, client_email, end_time)
        except myresource.MyResourceNotFound as e: # translate the resource manager exceptions to GENI exceptions
            raise geni_ex.GENIv3SearchFailedError("The desired my_resource(s) could not be found.")

        return self.lxml_to_string("<xml>omitted</xml>"), {'status' : '...omitted...'} # return the required results

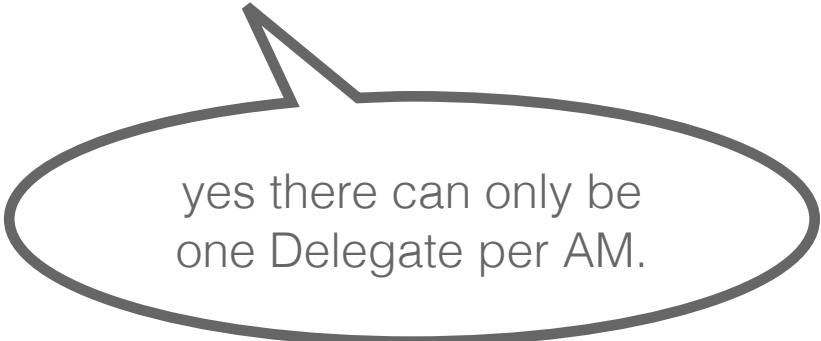
def setup():
    delegate = MyGENI3Delegate()
    handler = pm.getService('geniv3handler')
    handler.setDelegate(delegate)
```

needed knowledge



Delegate tasks

- Translate GENI API into Resource Manager(s) methods
- Translate the RSpecs into Resource Manager values (and back).
- Catch Resource Manager errors and re-throw as GENIv3....
- Translate the namespace from GENI to RM (e.g. URN ↔ UUIDs).
- Specify the needed privileges for authorization.
- De-multiplex to dispatch to different Resource Managers
(if you have multiple resource types in one AM).



yes there can only be
one Delegate per AM.

RM tasks

- Instantiate resources
- Manage persistence of reservations and resource state
- Check policies
- Avoid collisions of resources reservations /
Manage availability
- Throw domain-specific errors

more info

- Please see the [↻ wiki](#) for
 - Authentication / Authorization tools
 - RSpec generation assistance
 - More detailed description
- Checkout the code and look at the DHCP AM example
 - plugin: `dhcprm`
 - plugin: `dhcpgeni3`
 - API description of `geniv3rpc`

a table for two please

See what kind of bookings for resources are there and what is supported by AMsoil...

ways to schedule

There are two common types of scheduling

	best-effort	pre-booking
experimenter process	try and fail	convenient planning
scheduling constraints	current status only	current and future
data to maintain	past, current	past, current, future
resource usage pattern	typically sharing	typically exclusive use

types of resources

There are two different cardinalities for resource types.

	bounded	unbounded
available resources	limited	unlimited
availability check	boolean check	always available (possibly limited by the total load of booked resources)
resources identifiers	well known, limited number	non-clashing, possibly infinite

schedule API

We see different schedules, simple creation, bounded and unbounded.

```
import uuid
import amsoil.core.pluginmanager as pm

Schedule = pm.getService('schedule')
ip_schedule = Schedule("IPLease", 100) # create a schedule for IPs
vm_schedule = Schedule("VM", 100) # create a distinct schedule object for VMs

# create bounded reservations with dedicated resource ids
ip1 = ip_schedule.reserve(resource_id='192.168.1.1') # with mostly default values
ip2 = ip_schedule.reserve(resource_id='192.168.1.2')
# create a unbounded reservation
vm1 = vm_schedule.reserve(resource_id=str(uuid.uuid4()))

print len(ip_schedule.find()) # -> 2 (192.168.1.1, 192.168.1.2)
print len(vm_schedule.find()) # -> 1 (ec1f33f0-8443-11e3-baa7-0800200c9a66)
```

schedule API

We see complex reservation pre-booking and best-effort.

```
# complex creation for best effort (starts now)
ip1 = ip_schedule.reserve(
    resource_id='192.168.1.2',
    resource_spec={"additional_information" : [1,2,3] },
    slice_id='pizza',
    user_id='tom',
    start_time=datetime.utcnow(),
    end_time=datetime.utcnow() + timedelta(0,0,10,0))

# creation pre-booking with a default duration (from schedule constructor)
ip2 = ip_schedule.reserve(
    resource_id='192.168.1.3',
    start_time=datetime.utcnow() + timedelta(10,0,0,0)) # start in 10 days
```

schedule API

What a pickle! Where can I put my resource specific information?

there!

```
# complex creation for best effort (starts now)
ip1 = ip_schedule.reserve(
    resource_id='192.168.1.2',
    resource_spec={ "additional_information" : [1,2,3] },
    slice_id='pizza',
    user_id='tom',
    start_time=datetime.utcnow(),
    end_time=datetime.utcnow() + timedelta(0,0,10,0))
```

You can add custom info to each reservation (any [pickle](#)-able object).
If you can connect all info with reservations, no extra database needed.

hands on tips

Let's see how we can make our life even easier.

testing

- ✓ Fire up the Clearinghouse
- ✓ Start the AMsoil server
- ✓ Run omni to send a request
 - ✓ Check AMsoil's logs

```
gcf#      python src/gcf-ch.py
amsoil#    python src/main.py
amsoil#    tail -f log/amsoil.log
gcf#      python src/omni.py -o -a https://localhost:8001 -V 3 getversion
```

development mode

- Use the configuration tool to set `flask.debug = True`
 - Now the server [reloads its files every time](#) you change a file.

!! Careful: The client's certificate is now read from a pre-configured file.
- For debugging
 - Throw exceptions or
 - Write to the log to see what's going on.

logging

anywhere.py

```
import amsoil.core.log
logger=amsoil.core.log.getLogger('pluginname')
# logger is a decorated instance of Python's logging.Logger, so we only get one instance per name.

def somemethod():
    logger.info("doing really cool stuff...")
    logger.warn("Oh Oh...")
    logger.error("Ba-Boooom!!!")
```


configuration

anywhere.py

```
import amsoil.core.pluginmanager as pm
config = pm.getService("config")      # get the service
myvalue = config.get("mygroup.mykey") # retrieve a value
config.set("mygroup.mykey", myvalue)  # set a value
```

plugin.py

```
import amsoil.core.pluginmanager as pm
def setup():
    config = pm.getService("config") # get the service
    config.install("mygroup.mykey", "somedefault", "Some super description.") # install a config item
```



Always install the config keys and defaults on the plugin's setup method (install will not re-create/overwrite existing entries).

worker

The worker enables dispatching jobs to an external process (e.g. to perform longer tasks without blocking the client's request response).

anywhere.py

```
worker = pm.getService('worker') # get the service
worker.add("myservice", "mymethod", "parameter1") # run as soon as possible
worker.addAsRecurring("myservice", "mymethod", [1,2,3], 60) # run every minute
worker.addAsScheduled("myservice", "mymethod", None, datetime.now() + timedelta(0, 60*60*2)) # run in 2 hours
```

fire up the server (needs reboot when changing code)

```
amsoil# python src/main.py --worker
```

mailer

The mailer enables sending of plain-text mails.

anywhere.py

```
MailerClass = pm.getService('mailer')  
mailer = MailerClass('root@example.org', 'mail.example.org')  
mailer.sendMail("to@example.org", "Some Subject", "Some Body.")
```

! Delivering mail takes time.

!! Do not block the client's request handling too long.

✓ If you want to send multiple mails,
dispatch the delivery of mails to the [worker](#).

persistence


SQLAlchemy tutorial
7900 words

VS.

Need to know
926 words



you know it all



clone the repository

🔗 <https://github.com/fp7-ofelia/AMsoil.git>

then **read** the wiki

🔗 <https://github.com/fp7-ofelia/AMsoil/wiki>